

Semantic Ifs over Images: Logit-Readout Decisions with a Vision–Language Model

Samuel Otero Agraso

<https://github.com/XxSamaxX/discern>

Abstract

Recent open implementations read a decision directly from a language model’s answer-token logits instead of decoding text, turning a model into a runtime-defined conditional—a *semantic if*. We port this readout unchanged to a vision–language model and evaluate it on MMBench (en/dev, 4329 items) plus a small hand-built probe. Qwen3-VL-4B reaches 87.4% with one forward pass per item at 91 ms, and 82.2% under MMBench’s CircularEval protocol, which rotates the answer options.

Our main results concern the confidence signal rather than the score. The softmax over answer slots, which these systems report as a probability, is ≥ 0.99 on 52% of the model’s own errors: it claims near-certainty on more than half of what it gets wrong. The logit gap behind that same number predicts correctness with AUC 0.894 and supports a usable operating point—thresholding at 16 nats answers 66% of items at 98.4% accuracy, against an 87.4% base rate. The gap also predicts order sensitivity almost perfectly (AUC 0.980): rotating options changes the winner on 11.7% of items, and those items have a median gap of 2.7 nats against 22.8 for the stable ones, which bounds when order calibration is worth its n forward passes.

On POPE (9000 binary presence questions) the same readout scores 89.6%, but the gap’s AUC falls to 0.777 and the probability is ≥ 0.99 on 86% of the errors: the signal weakens as the option set shrinks, and a threshold tuned on four-way questions does not transfer to two-way ones.

Two negative results survive the scale-up. First, difficulty is not epistemics: on 12 normative questions no model can ground, the gap distribution shifts (median 8.8 vs 20.9 nats, AUC 0.788, $p = 0.006$) but overlaps so heavily that the best threshold still admits 7 of 12—the model asserts that a photograph requires legal consent at 15.6 nats, as confidently as it reports that no dog is present at 26.4. Second, our own 22-item probe scored 22/22 with every softmax probability above 0.9993, which we initially read as universal saturation; the 4329-item run shows that this was a property of an over-easy benchmark, and we report it as a caution about small author-written test sets.

1 Introduction

A recurring pattern in recent agent tooling is to stop asking a language model for text and ask it for a decision. The construction is simple: phrase the decision as a lettered multiple choice, run a single forward pass, and read the logits at the final position restricted to the tokens A, B, C. There is no decoding loop, so the cost is one prefill. The open SemIf project [3] publishes this readout as a base-

line, and the JEV-CPU port [6] demonstrates it on a laptop CPU with a 0.6B model. Practitioners describe the result as a *semantic if*: a conditional whose predicate is written in natural language at runtime rather than compiled.

The attraction for computer vision is immediate. A detector trained on a fixed label set answers “is there a person” but cannot answer “is this photograph posed”, “is the white balance wrong”, or “is the meal already in progress”. These are not classes; they are predicates. If the readout transfers to a vision–language model, the predicate becomes free text and no training set is required.

This paper reports what happens when it does transfer. Our contributions are:

1. A port of the readout to a vision–language model in which the scoring path is byte-identical to the text implementation; only prompt construction and model loading differ (Section 3).
2. An evaluation on MMBench en/dev: 87.4% at 91 ms per item, 82.2% under CircularEval (Section 4.5), plus a 22-item adversarial probe the model answers in full (Section 4.1).
3. The reported probability is not merely miscalibrated but actively misleading: it exceeds 0.99 on 52% of the model’s errors. The logit gap behind it predicts correctness with AUC 0.894 and supports abstention at 98.4% accuracy over 66% of items (Sections 4.2 and 4.5).
4. A characterisation of position bias: 1–2 nats in magnitude, flipping the winner on 11.7% of items, with the gap predicting order-stability at AUC 0.980. This bounds when order calibration earns its n forward passes (Section 4.3).
5. A negative result we consider the most consequential for deployment: the gap does not separate answerable perceptual questions from ungroundable normative ones well enough to gate them (Section 4.4).
6. A methodological caution, learned the hard way: our own 22-item benchmark reported universal softmax saturation and zero order sensitivity, and the 4329-item run refuted both.

2 Related work

Open-vocabulary detection attacks the same limitation from the other side. CLIP-style image–text alignment [8] underpins detectors such as OWL-ViT [7], Grounding DINO [5] and YOLO-World [1], which accept a free-text

vocabulary and localise its members. These systems remain far faster than a vision–language forward pass and they return boxes, which the readout studied here does not. They are, however, restricted to naming things: they cannot evaluate a predicate over a scene. The two approaches are complementary rather than competing, and a natural composition is to let a detector propose regions and a semantic if evaluate predicates on the crops.

Our negative results connect to two literatures. Position bias in multiple choice prompting is well documented [11], and contextual calibration [10] is the standard remedy; we measure the bias in nats and find it an order of magnitude smaller than the decision margins on perceptual items, which bounds when that remedy is worth its cost. Separately, modern networks are known to be poorly calibrated [2]; we show a sharper failure specific to this readout, where the reported probability is not merely miscalibrated but constant. Finally, object hallucination benchmarks such as POPE [4] establish that vision–language models assert the presence of absent objects; our false-premise and absence items probe the same behaviour and, notably, do not reproduce it at this model scale.

3 Materials and methods

3.1 The decision readout

A decision is a tuple of evidence, a criterion, and 2–16 options. The prompt is a JSON payload naming each option with a letter, under a system instruction to answer with that letter alone. The implementation verifies that each letter encodes to exactly one token with an exact decode round trip, and that appending the letter does not alter the tokenisation of the prompt that precedes it; both checks raise rather than silently degrade. A single forward pass with `use_cache=False` yields the final-position logit vector, which is indexed at the answer-slot tokens. A softmax over those few entries produces the reported probabilities.

3.2 Extension to images

We changed exactly two things. The user turn carries an image part before the JSON payload, and the model is loaded as an image-text-to-text model in bfloat16 on the GPU. The slot verification, the forward call and the logit indexing are unchanged from the text implementation, so results are directly comparable. At the processor’s default resolution a 640×480 photograph expands to roughly 300 visual tokens, giving prompts of 351–413 tokens in our benchmark.

3.3 Benchmark

We wrote 22 decisions over four photographs from COCO *val2017*. Twelve are ordinary perceptual judgements, three of which a fixed-vocabulary detector could also answer and act as controls. Ten target families known to be difficult for vision–language models: exact counting, absence, a false-premise question (the colour of a car in an image containing no car), an unanswerable question (the time of day in

a windowless interior), fine-grained comparison, and spatial relation. Ground truth was written by the author from the images. One label was initially wrong—we judged a brand mark illegible, and the model read it correctly—and was corrected after visual inspection of an enlarged crop; we report this because the corrected item is the hardest in the set.

A second set of 12 *normative* questions was written to be structurally identical but institutionally ungrounded: whether consent must be obtained, whether a retention policy requires deletion, whether a model release has expired. Each carries one sentence of evidence and two options. No model can answer these from the evidence given, because the answer depends on a policy it has never seen.

3.4 Order calibration and the logit gap

For each decision we optionally re-score under all cyclic rotations of the option list and average in logit space. We record, per decision, the gap in nats between the two highest mean logits, whether the winner changes across rotations, and the mean logit received by each letter slot, which with rotation is a clean estimate of position bias.

3.5 Hardware and models

All experiments run on one machine: an Intel i9-12900F (8 performance and 8 efficiency cores) with 31 GiB of RAM and an NVIDIA RTX 3080 with 12 GiB. The vision model is Qwen3-VL-4B-Instruct in bfloat16 on the GPU. The text models are Qwen3-0.6B [9] and Qwen3.5-4B in float32 on the CPU; the latter is the model the upstream manifest [3] designates for this readout, and the CPU port substitutes the former.

4 Results

4.1 Perceptual accuracy

Qwen3-VL-4B answers 22 of 22 correctly. All four adversarial families are handled: it counts two cats and two remote controls, denies the presence of a dog and of a cat in the rooms that contain neither, answers “there is no car” to the false-premise question at a gap of 22.1 nats, and answers “cannot be determined” to the unanswerable one at 15.5 nats. It reads an Adidas mark on a cap occupying roughly 20×15 pixels of a 640×480 image. The two lowest gaps in the set are the fine-grained comparison (7.3 nats) and the spatial relation (8.1 nats), which are the two items we judged hardest a priori. Throughput is 110–140 ms per decision at 8.4 GiB of VRAM, after a first forward pass of 11.2 s attributable to CUDA warm-up.

4.2 The reported probability is constant

Every one of the 22 decisions reports a winning-option probability of at least 0.9993 (Figure 1, top). Taken at face value the system is certain about all of them, and no threshold on that number can separate the item it found hardest




Figure 1: The same 22 decisions and the same forward passes, read two ways. Top: the softmax probability of the winning option is indistinguishable from 1.0 everywhere. Bottom: the logit gap behind those probabilities spans 7.3–26.4 nats and orders the decisions by difficulty.

from the item it found easiest. On MMBench the saturation is severe but not total—the minimum is 0.325 and 88% of items exceed 0.99—so this figure overstates the effect; what survives at scale is the consequence, that the probability is ≥ 0.99 on 52% of the model’s errors (Section 4.5). The information is not absent from the model; it is destroyed by the readout. A softmax over two or three slots maps any gap beyond roughly 7 nats to a probability that rounds to one. The gaps themselves (Figure 1, bottom) are well spread and monotone in our a priori difficulty ordering. The upstream implementation labels its output “*conditional option score; uncalibrated as decision confidence*”, which is correct but understates the problem: the quantity is not a weak confidence signal, it is very nearly a constant. Practitioners who need a threshold should apply it to the gap in nats.

4.3 Order sensitivity is small

Averaged over rotations, the mean logit assigned to each letter slot differs by 0.2–2.4 nats, with no consistent preference for the first slot. Against perceptual gaps of 7.3 nats and above, this is an order of magnitude too small to matter, and indeed the winner changes under rotation in 0 of 22 perceptual decisions. Order calibration on this workload costs n forward passes and buys nothing. It is not free elsewhere: in the text-model experiments of Section 4.7, where gaps fall to 0.8–2.3 nats, the same 1–2 nats of position bias decides the answer outright and the winner flips. A practical rule follows: rotate and average only when the single-pass gap is below roughly 3 nats.




Figure 2: Logit gaps for 22 perceptual decisions and 12 normative ones under the same model and readout. The distributions differ (AUC 0.788, $p = 0.006$) but overlap over most of their range; the shaded band marks the overlap.

Table 1: Qwen3-VL-4B on MMBench en/dev (4329 items). CircularEval requires the correct answer under every rotation of the options.

Protocol	Accuracy	ms/item	AUC gap
VanillaEval	87.4%	91	0.894
CircularEval	82.2%	347	0.956

4.4 The gap does not gate normative questions

If the gap tracks difficulty within a family, the natural hope is that it also flags questions the model has no business answering. We tested this on the 12 normative items, which are unanswerable in principle rather than merely hard.

The hope fails (Figure 2). The distributions do differ: median 8.8 nats for normative against 20.9 for perceptual, giving an AUC of 0.788 (Mann–Whitney $U = 208$, $z = 2.74$, $p = 0.006$). But 7 of 12 normative items fall inside the perceptual range, and the best achievable threshold, 7.3 nats, retains all 22 perceptual decisions only by admitting 7 of the 12 normative ones—79% accuracy for a filter whose entire purpose is to catch the remainder. Individual cases are worse than the aggregate suggests. The model asserts that a photograph with a colour cast should enter the correction queue at 23.1 nats and that publication must be held for consent at 15.6, both policy questions, at gaps indistinguishable from its report that there is no dog in the picture (26.4). Only 1 of 12 normative items flips under rotation.

Two of our normative items are arguably entailed by their evidence, which inflates their gaps and weakens the contrast; we note this rather than remove them post hoc.

4.5 Validation at scale on MMBench

Our own benchmark returned 22/22, so it could not measure whether the gap *predicts* anything—there were no errors to predict. MMBench supplies 547 of them under VanillaEval.

The gap separates correct from incorrect answers with AUC 0.894 ($z = 29.9$), with a median of 22.4 nats when the model is right against 5.2 when it is wrong. Thresholding on it yields a genuine abstention policy: at 12 nats

Table 2: POPE, 9000 binary presence questions, 1500 yes and 1500 no per split. Recall is identical across splits because the positive half is shared; only the distractors change.

Split	Acc.	Prec.	Rec.	Yes-rate
random	91.2%	98.2%	83.9%	42.7%
popular	89.3%	94.0%	83.9%	44.6%
adversarial	88.3%	92.0%	83.9%	45.6%
all	89.6%	94.7%	83.9%	44.3%

the system answers 75% of items at 96.8% accuracy, and at 16 nats 66% at 98.4%, against an 87.4% base rate. The reported probability supports nothing comparable—it is at least 0.99 on 88% of all items and, more to the point, on 52% of the errors.

CircularEval costs 5.2 points (87.4% to 82.2%) because the winner changes under rotation on 507 of 4329 items, 11.7%. Those items are not randomly distributed: their median gap is 2.7 nats against 22.8 for the order-stable ones, and the gap predicts stability with AUC 0.980. The rule we proposed in Section 4.3 from a 22-item probe—rotate only below roughly 3 nats—was written before this dataset was downloaded and lands on the empirical median of the items that actually flip.

The weakest categories are `spatial_relationship` (34.5% under CircularEval, against a 25% chance rate) and `image_quality` (57.3%). The two lowest gaps in our hand-built probe were its spatial-relation item (8.1 nats) and its white-balance item (7.8). A confidence signal read off a handful of examples identified the categories that collapse at scale.

4.6 Binary questions: the signal weakens

MMBench items carry four options. POPE [4] carries two, which is the worst case for a readout that softmaxes over answer slots, and it probes exactly the family—object presence and absence—on which our hand-built probe was perfect.

Accuracy is 89.6% with an F1 of 89.0 (Table 2), but the confidence signal degrades sharply. The gap separates correct from incorrect answers at AUC 0.777, against 0.894 on MMBench, and the reported probability is ≥ 0.99 on 86% of the errors rather than 52%. The median gap when the model is wrong rises from 5.2 nats on MMBench to 16.4 here: with two slots the model is not merely wrong, it is confidently wrong. The practical consequence is that a threshold tuned on multi-way questions does not transfer to binary ones, and that a semantic if should be given three or more options wherever the decision admits them.

Order sensitivity moves the other way, and consistently with our account of it: only 0.6% of items change winner under rotation, against 11.7% on MMBench, because binary gaps are large. Those that do flip have a median gap of 0.8 nats, and the gap predicts stability at AUC 0.998.

POPE also shows where object hallucination hides. Its headline metric, the yes-rate, is 42.7–45.6%—below the 50% base rate—so on that metric the model looks clean, and our own prediction that it would over-assert presence

Table 3: Latency of one decision, by model and placement. CPU figures are float32; thread counts are the best of a sweep.

Model	Device	Threads	ms/forward
Qwen3-VL-4B	RTX 3080	—	110–140
Qwen3-0.6B	i9-12900F	20	340–450
Qwen3.5-4B	i9-12900F	8	1359
Qwen3.5-4B	i9-12900F	20	2392

was wrong. The effect is real but sits in precision: 98.2% to 94.0% to 92.0% across the random, popular and adversarial splits, while recall stays at exactly 83.9%. Co-occurring distractors do pull the model into false positives; an aggregate yes-rate averages that away.

4.7 Hybrid CPU–GPU execution

Because the two devices are idle at different times, we built a two-stage cascade: the GPU resolves visual predicates over an image, and a CPU model decides an action from the resulting verdicts without ever seeing pixels. Run as two threads with a queue between them, the stages overlap—PyTorch releases the interpreter lock during both CUDA and CPU forwards—and wall-clock time falls from 3.49 s to 2.72 s over four images, a 22% saving.

The thread sweep in Table 3 is a practical result in its own right. On this hybrid CPU, restricting inference to the 8 performance cores is $1.9\times$ faster than using 20 threads; scheduling work onto the efficiency cores actively hurts. Any measurement taken at the default thread count understates CPU inference on such parts, and our own first cascade measurements did.

One optimisation that looks free is not. Questions about one image share a prefix—the system turn and roughly 300 visual tokens—so caching it and running only each question’s tail should save most of the work. The cached path is numerically exact, reproducing every probability to the last bit, and it is 13–14% slower, flat across images of 384, 852 and 1812 tokens with no crossover: it loses the fast attention kernel, which costs more than recomputing the prefix. Image preprocessing accounts for 3% of the time and the forward for 97%.

The policy stage itself did not work, and the way it failed is instructive. Our first design asked one four-option question whose options were not mutually exclusive—an image can need both colour correction and a consent hold. The model oscillated between two *correct* answers, which is a defect of the question, not of the model. Restructuring the stage as independent binary predicates composed in ordinary Python removed the ambiguity. Even then, Qwen3-0.6B answered 1 of 4 policy cases correctly and was unstable under rotation in 3; Qwen3.5-4B answered 3 of 4 and was stable in all 4, but at gaps of 0.7–2.3 nats, below the threshold at which we would trust any of them. The larger model also shifts the pipeline from balanced to 12:1 against the CPU stage, which removes the reason to split the work in the first place.

5 Discussion

The perceptual result is strong and the systems result is modest; the two negative results are what we would want a practitioner to take away.

The first is a readout defect with a one-line fix. Reporting a softmax over two or three answer slots produces a number that is always approximately one, and pipelines that threshold on it are thresholding on noise. The gap in nats is already computed and should be reported instead.

The second is harder to fix and easier to get wrong in production. A semantic if will answer a question about your organisation’s consent policy with the same apparent conviction it brings to counting cats, and no signal we measured reliably distinguishes the two. The pattern is therefore safe where the predicate is perceptual and grounded in the input, and unsafe where the predicate is normative. Business policy belongs in code. This is also the lesson of the cascade: the mechanism is called a semantic *if* for a reason, and composing several binary ifs in a host language is both more accurate and more debuggable than asking one model to resolve a multi-way policy.

Limitations. One model, one prompt template, one language, a single machine and a single run, without repetition or confidence intervals on latency. MMBench dev is public and may appear in the model’s training data, so 87.4% is not evidence about generalisation; the claims we defend are about the confidence signal, which is measured within the same run and does not depend on the absolute score. Our hand-built probe is small and author-annotated, one of its labels was demonstrably wrong before correction, and its 22/22 misled us about the extent of softmax saturation until the larger run corrected it—which is itself the clearest argument in this paper against trusting a small author-written benchmark. The normative set has 12 items, two of which are arguably entailed by their evidence.

6 Conclusions

Reading a decision from answer-token logits transfers to vision-language models without modification and is accurate and fast on grounded perceptual predicates: 22 of 22 at 110–140 ms per decision on one consumer GPU, robust to option order, and correct on counting, absence, false premises and unanswerable questions. The probability such systems report is not usable as confidence—it exceeded 0.9993 on every decision we ran—while the logit gap behind it is a good within-family difficulty signal. That signal does not generalise across question families: it fails to separate perceptual questions from normative ones well enough to gate them, and the model is confidently wrong about policy. The signal is also not scale-free: halving the option set from four to two drops its AUC from 0.894 to 0.777 and leaves the model confidently wrong, so a threshold must be set per question shape rather than once. Use semantic ifs for perception, give them three or more options where the decision allows, keep policy in code, and threshold on nats rather than on probabilities.

Code and data

The method is packaged as `discern`, a one-call API in which a below-threshold verdict refuses to be used as a boolean condition. Code, benchmark, raw results and this paper’s sources are available at <https://github.com/XxSamaxX/discern>. The scoring engine is vendored from SemIf [3] under the MIT licence.

References

- [1] Tianheng Cheng, Lin Song, Yixiao Ge, Wenyu Liu, Xinggang Wang, and Ying Shan. YOLO-World: Real-time open-vocabulary object detection. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [2] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *International Conference on Machine Learning (ICML)*, 2017.
- [3] Theo C. J. Lee. SemIf: Open baselines for runtime-defined semantic decisions, 2025. <https://github.com/TheoLeeCJ/SemIf>.
- [4] Yifan Li, Yifan Du, Kun Zhou, Jinpeng Wang, Wayne Xin Zhao, and Ji-Rong Wen. Evaluating object hallucination in large vision-language models. In *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- [5] Shilong Liu et al. Grounding DINO: Marrying DINO with grounded pre-training for open-set object detection. *arXiv preprint arXiv:2303.05499*, 2023.
- [6] Meanblock. JEV-CPU: Semantic ifs from open models on a laptop CPU, 2026. <https://huggingface.co/Meanblock/JEV-CPU>.
- [7] Matthias Minderer et al. Simple open-vocabulary object detection with vision transformers. In *European Conference on Computer Vision (ECCV)*, 2022.
- [8] Alec Radford et al. Learning transferable visual models from natural language supervision. In *International Conference on Machine Learning (ICML)*, 2021.
- [9] An Yang et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [10] Tony Z. Zhao, Eric Wallace, Shi Feng, Dan Klein, and Sameer Singh. Calibrate before use: Improving few-shot performance of language models. In *International Conference on Machine Learning (ICML)*, 2021.
- [11] Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *International Conference on Learning Representations (ICLR)*, 2024.